

# Machine learning

## Language models

Katya Artemova

CS HSE

February 25, 2019

# Overview

- 1 Language modeling task
- 2 Count-based language models
- 3 Neural language models
  - Recurrent neural network
  - Gated architectures
  - RNN generators
- 4 Bonus: RecNN

# Language model

- 1 Compute the probability of a sequence of words:

$$P(w_1, w_2, \dots, w_n)$$

- 2 Predict next word:

$$P(w_n | w_1, w_2, \dots, w_{n-1})$$

LMs help are used to:

- Machine translation: choose best translation
- Spell checking: find incorrect word
- Speech recognition: choose best transcription
- Predict next word in your smartphone
- Generate poems, summaries, answers, etc.

# Overview

- 1 Language modeling task
- 2 Count-based language models
- 3 Neural language models
  - Recurrent neural network
  - Gated architectures
  - RNN generators
- 4 Bonus: RecNN



A. A. Mapcon (1886).

# Markov assumptions

- ① Chain rule:

$$P(W) = P(w_1, w_2, \dots, w_n) = \prod_i P(w_i | w_1, \dots, w_{i-1})$$

- ② Maximum likelihood estimates of probabilities:

$$P(w_i | w_1, \dots, w_{i-1}) = \frac{\text{count}(w_1, w_2, \dots, w_i)}{\text{count}(w_1, \dots, w_{i-1})}$$

- ③ Markov assumption ( $k$ -th order):

$$P(w_i | w_1, \dots, w_{i-1}) \approx P(w_i | w_{i-k}, \dots, w_{i-1})$$

# $n$ -gram models

- ① Unigram models:  $P(W) = P(w_1, \dots, w_n) \approx \prod_i P(w_i)$
- ② Bigram models:  $P(W) = P(w_1, \dots, w_n) \approx \prod_i P(w_i | w_{i-1})$
- ③ Perplexity:  $PP(W) = P(w_1, w_2, \dots, w_n)^{-\frac{1}{N}}$   
The lower perplexity, the better the model predicts an unseen test
- ④ Smoothing:  $P(w_i | w_1, \dots, w_{i-1}) = \frac{\text{count}(w_1, w_2, \dots, w_i) + 1}{\text{count}(w_1, \dots, w_{i-1}) + \alpha |V|}$ , where  $|V|$  is the size of dictionary
- ⑤ Interpolation:  $\hat{P}(w_i | w_{i-1}) = \lambda P_{MLE}(w_i | w_{i-1}) + (1 - \lambda) P_{MLE}(w_i)$

# $n$ -gram models for text generation

Given  $w_i$ :

- 1 choose the next most probable  $w_{i+1}$
- 2 randomly select sample from this probability distribution of next words



Ветхий Алгоритм  
@alg\_testament

Follow

дети не должны использоваться эти две функции обращения к массиву environ.

6:53 AM - 21 Aug 2018



Ветхий Алгоритм  
@alg\_testament

Follow

все проклятия, написанные в книге, поддерживают 32- и 64-разрядные смещения

[https://twitter.com/alg\\_testament](https://twitter.com/alg_testament)



# $n$ -gram models for IR

Given documents  $D$  and query  $q$ , estimate the probability of generating the query text from a document language model:

- Rank documents by the probability that the query could be generated by the document model;
- Calculate  $P(d|q)$  to rank the documents:  $P(d|q) \propto P(q|d)P(d)$
- Assuming prior is uniform, unigram model:  $P(q|d) = \prod_i P(q_i|d)$
- MLE:  $P(q_i|d) = \frac{\text{count}(q_i, d)}{|d|}$

# Reading

- 1 Stuart Russell, Peter Norvig. Artificial Intelligence: A Modern Approach, Ch. 15
- 2 Dan Jurafsky, James H. Martin. Speech and Language Processing, Ch. 3, Ch. 8

# Sequential data

- ① Time series
  - ▶ Financial data analysis: stock market, commodities, Forex
  - ▶ Healthcare: pulse rate, sugar level (from medical equipment and wearables)
- ② Text and speech: speech understanding, text generation
- ③ Spatiotemporal data
  - ▶ Self-driving and object tracking
  - ▶ Plate tectonic activity
- ④ Physics: jet identification
- ⑤ etc.

# Sequence modelling I

## Sequence labelling

- ①  $\mathbf{x} = x_1, x_2, \dots, x_n, x_i \in V$ , - objects
- ②  $\mathbf{y} = y_1, y_2, \dots, y_n, y_i \in \{1, \dots, L\}$  - labels
- ③  $\{(\mathbf{x}^{(1)}, \mathbf{y}^{(1)}), (\mathbf{x}^{(2)}, \mathbf{y}^{(2)}), \dots, (\mathbf{x}^{(m)}, \mathbf{y}^{(m)})\}$  – training data
- ④ exponential number of possible solutions : if  $\text{length}(\mathbf{x}) = n$ , there are  $L^n$  possible solutions

Classification problem:  $\gamma : \mathbf{x} \rightarrow \mathbf{y}$

- ① Speech recognition:  $x$  – spoken words,  $y$  – transcription
- ② Genome annotation:  $x$  – DNA,  $y$  – genes

# Sequence modelling II

## Sequence classification

- ①  $\mathbf{x} = x_1, x_2, \dots, x_n, x_i \in V$ , - objects
- ②  $y \in \{1, \dots, L\}$  - labels
- ③  $\{(\mathbf{x}^{(1)}, y_1), (\mathbf{x}^{(2)}, y_2), \dots, (\mathbf{x}^{(m)}, y_m)\}$  - training data

Classification problem:  $\gamma : \mathbf{x} \rightarrow y$

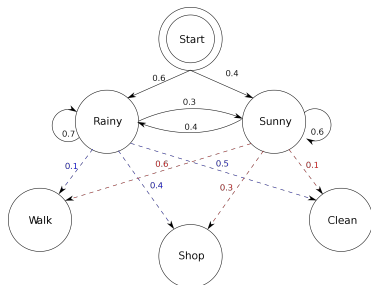
- ① Activity recognition:  $x$  - pulse rate,  $y$  - activity (walking, running, peace)
- ② Opinion mining:  $x$  - sentence,  $y$  - sentiment (positive, negative)
- ③ Trading:  $x$  - stock market,  $y$  - action (sell, buy, do nothing)

# Traditional ML approaches to sequence modelling

- Hidden Markov Models (HMM)
- Conditional Random Fields (CRF)
- Local classifier: for each  $x$  define features, based on  $x_{-1}$ ,  $x_{+1}$ , etc, and perform classification  $n$  times

Problems:


- 1 Markov assumption: fixed length history
- 2 Computation complexity



- 1 Language modeling task
- 2 Count-based language models
- 3 Neural language models
  - Recurrent neural network
  - Gated architectures
  - RNN generators
- 4 Bonus: RecNN

# Recurrent neural network

- Input: sequence of vectors
- $x_{1:n} = x_1, x_2, \dots, x_n, x_i \in \mathbb{R}^{d_{in}}$
- Output: a single vector  
 $y_n = RNN(x_{1:n}), y_n \in \mathbb{R}^{d_{out}}$
- For each prefix  $x_{1:i}$  define an output vector  $y_i$ :  
 $y_i = RNN(x_{1:i})$
- $RNN^*$  is a function returning this sequence for input sequence  $x_{1:n}$ :  
 $y_{1:n} = RNN^*(x_{1:n}), y_i \in \mathbb{R}^{d_{out}}$





# Sequence modelling with RNN

## 1 Sequence labelling

Produce an output  $y_i$  for each input RNN reads in. Put a dense layer on top of each output to predict the desired class of the input

$$p(l_j | \mathbf{x}_j) = \text{softmax}(RNN(\mathbf{x}_{1:j}) \times W + b)_{[j]}$$


## 2 Sequence classification

Put a dense layer on top of RNN to predict the desired class of the sequence after the whole sequence is processed

$$p(l_j | \mathbf{x}_{1:n}) = \text{softmax}(RNN(\mathbf{x}_{1:n}) \times W + b)_{[j]}$$


# More details on RNN

- $RNN^*(x_{1:n}, s_0) = y_{1:n}$
- $y_i = O(s_i)$  – simple activation function
- $s_i = R(s_{i-1}, x_i)$ , where  $R$  is a recursive function,  $s_i$  is a state vector
- $s_0$  is initialized randomly or is a zero vector
- $x_i \in \mathbb{R}^{d_{in}}$ ,  $y_i \in \mathbb{R}^{d_{out}}$ ,  $s_i \in \mathbb{R}^{f(d_{out})}$
- $\theta$  – shared weights




# More details on RNN

- $s_i = R(x_i, s_{i-1}) = g(s_{i-1}W^s + x_iW^x + b)$
- $y_i = O(s_i) = s_i$
- $y_i, s_i, b \in \mathbb{R}^{d_{out}}, x_i \in \mathbb{R}^{d_{in}}$
- $W^x \in \mathbb{R}^{d_{in} \times d_{out}}, W^s \in \mathbb{R}^{d_{out} \times d_{out}}$




# RNN unrolled



$$\begin{aligned} s_4 &= R(s_3, x_4) = R(R(s_2, x_3), x_4) = R(R(R(s_1, x_2), x_3), x_4) = \\ &= R(R(R(R(s_0, x_1), x_2), x_3), x_4) \end{aligned}$$

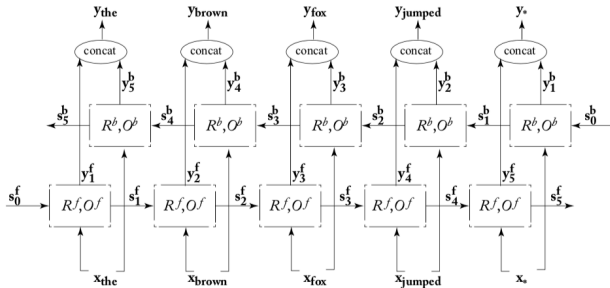
# Bidirectional RNN (Bi-RNN)

The input sequence can be read from left to right and from right to left.  
Which direction is better?



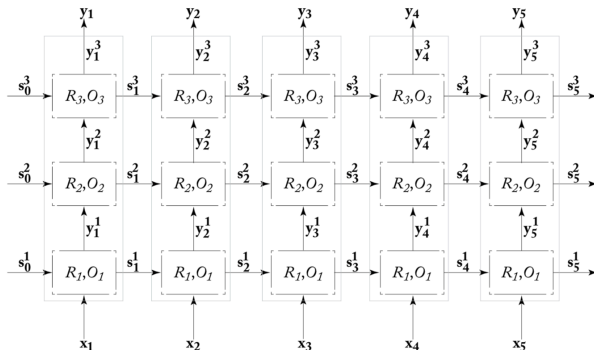
$$biRNN(x_{1:n}, i) = y_i = [RNN^f(x_{1:i}); RNN^r(x_{n:i})]$$

# Bi-RNN



$$biRNN^*(x_{1:n}, i) = y_{1:n} = biRNN(x_{1:n}, 1) \dots biRNN(x_{1:n}, n)$$


# Multilayer RNN



Connections between different layers are possible too:  $y^2_1 = \text{concat}(x_1, y^1_1)$

# Sequence labelling


- Output  $\hat{t}_i$  for each input  $x_{1,i}$
- Local loss:  $L_{local}(\hat{t}_i, t_i)$
- Global loss:  
$$L(\hat{t}_n, t_n) = \sum_i L_{local}(\hat{t}_i, t_i)$$
- $L$  can take any form: cross entropy, hinge, margin, etc.






# Sequence classification

- $\hat{y}_n = O(s_n)$
- prediction =  $MLP(\hat{y}_n)$
- Loss:  $L(\hat{y}_n, y_n)$
- $L$  can take any form: cross entropy, hinge, margin, etc.



# Backpropagation through time



$$s_i = R(x_i, s_{i-1}) = g(s_{i-1}W^s + x_iW^x + b)$$

$$\text{Chain rule: } \frac{\partial L}{\partial w} = \frac{\partial L}{\partial p(\hat{y}_5)} \frac{\partial p(\hat{y}_5)}{\partial s_4} \left( \frac{\partial s_4}{\partial w} + \frac{\partial s_4}{\partial s_3} \frac{\partial s_3}{\partial w} + \frac{\partial s_4}{\partial s_3} \frac{\partial s_3}{\partial s_2} \frac{\partial s_2}{\partial w} + \dots \right)$$

# Vanishing gradient problem

Chain rule:  $\frac{\partial L}{\partial w} = \frac{\partial L}{\partial p(\hat{y}_5)} \frac{\partial p(\hat{y}_5)}{\partial s_4} \left( \frac{\partial s_4}{\partial w} + \frac{\partial s_4}{\partial s_3} \frac{\partial s_3}{\partial w} + \frac{\partial s_4}{\partial s_3} \frac{\partial s_3}{\partial s_2} \frac{\partial s_2}{\partial w} + \dots \right)$   
 $g$  – sigmoid

- ❶ Many sigmoids near 0 and 1
  - ▶ Gradients  $\rightarrow 0$
  - ▶ Not training for long term dependencies
- ❷ Many sigmoids  $> 1$ 
  - ▶ Gradients  $\rightarrow +\infty$
  - ▶ Not training again


Solution: gated architectures (LSTM and GRU)

- 1 Language modeling task
- 2 Count-based language models
- 3 Neural language models
  - Recurrent neural network
  - Gated architectures
  - RNN generators
- 4 Bonus: RecNN

# Controlled memory access


- Entire memory vector is changed:  $s_{i+1} = R(x_i, s_i)$
- Controlled memory access:  $s_{i+1} = g \odot R(x_i, s_i) + (1 - g)s_i$   
 $g \in [0, 1]^d, s, x \in \mathbb{R}^d$
- Differential gates:  $\sigma(g), g' \in \mathbb{R}^d$
- This controllable gating mechanism is the basis of the LSTM and the GRU architectures

# Long short term memory



<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>


# Long short term memory



$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

# Long short term memory




$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>




# Long short term memory



$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

# Long short term memory




$$o_t = \sigma(W_o [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh(C_t)$$

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

# Gated recurrent unit



$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh(W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

- 1 Language modeling task
- 2 Count-based language models
- 3 Neural language models
  - Recurrent neural network
  - Gated architectures
  - RNN generators
- 4 Bonus: RecNN

# Sequence generation

- Teacher forcing:


$x := \langle s \rangle x, y := x \langle /s \rangle$

$x := \langle s \rangle x_1 x_2 \dots x_n$

$y := x_1 x_2 \dots x_n \langle /s \rangle$

- $\hat{t}_{j+1} \sim p(t_{j+1} = k | t_{1:j})$

- $p(t_{j+1} = k | t_{1:j}) = f(RNN(\hat{t}_{1:j}))$   
 $f(x) = \text{softmax}(MLP(x))$





# Sequence generation

- Examples of generated texts:

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

- Examples of generated MIDI music:

<https://towardsdatascience.com/>

[how-to-generate-music-using-a-lstm-neural-network-in-keras](#)

# Conclusion

Topics covered:

- 1 RNN is a powerful tool for sequence modeling
- 2 RNN usage scenarios: sequence labelling, sequence classification, sequence generation
- 3 RNN layers can be reversed  $\rightarrow$  bidirectional RNN
- 4 RNN layers can be stacked  $\rightarrow$  deep RNN
- 5 RNN suffers from gradient vanishing problem  $\rightarrow$  LSTM, GRU

Topics not covered:

- 1 seq2seq models
- 2 Attention mechanism in RNN
- 3 Recursive neural networks




- 1 Language modeling task
- 2 Count-based language models
- 3 Neural language models
  - Recurrent neural network
  - Gated architectures
  - RNN generators
- 4 Bonus: RecNN



# Modeling trees with Recursive NN

- Input:  $x_1, x_2, \dots, x_n$
- A binary tree  $T$  can be represented as a unique set of triplets  $(i, k, j)$ , s.t.  $i < k < j$ ,  $x_{i:j}$  is parent of  $x_{i:k}, x_{k+1:j}$
- RecNN takes as an input a binary tree and returns as output a corresponding set of inside state vectors  $s_{i:j}^A \in \mathbb{R}^d$
- Each state vector  $s_{i:j}^A$  represents the corresponding tree node  $q_{i:j}^A$  and encodes the entire structure rooted at that node

# RecNN

- Input:  $x_1, x_2, \dots, x_n$  and a binary tree  $T$
- $\text{RecNN}(x_1, x_2, \dots, x_n, T) = \{s_{i:j}^A \in \mathbb{R}^d \mid q_{i:j}^A \in T\}$
- $s_{i:i}^A = v(x_i)$
- $s_{i:j}^A = R(A, B, C, s_{i:k}^B, s_{k+1:j}^C),$   
 $q_{i:k}^B \in T, q_{k+1:j}^C \in T$
- $R(A, B, C, s_{i:k}^B, s_{k+1:j}^C) = g([s_{i:k}^B, s_{k+1:j}^C]W)$





# Reading

- 1 Yoav Goldberg. Neural Network Methods for Natural Language Processing, Ch. 14-19
- 2 Ian Goodfellowm Yoshua Bengio, Aaron Courville. Deel learning, Ch. 10

# References I



Richard Socher, Alex Perelygin, Jean Wu, Jason Chuang, Christopher D Manning, Andrew Ng, and Christopher Potts, *Recursive deep models for semantic compositionality over a sentiment treebank*, Proceedings of the 2013 conference on empirical methods in natural language processing, 2013, pp. 1631–1642.